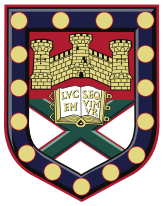

University of Exeter

Can statistics help us to understand deep learning?

Johannes Smit

May 2019

Abstract

All theses/dissertations must include an abstract of approximately 300 words bound in with each copy and placed so as to follow the title page.

Acknowledgements

Hi mum.

Contents

Chapter 1

Introduction

This is chapter one.

- What is machine learning?
- Machine learning is very complex to human understanding
- Machine learning is widely used in many high stakes scenarios — give examples
- It is important to be able to look inside the 'black box' of machine learning — explain how this would be useful in each of the given examples
- We can use statistical methods to try and recover the information in a ML algorithm — give some idea how

Chapter 2

Literature Review

This is chapter two.

- Machine Learning
  - Neural Networks
  - Deep Learning
- Stepwise Regression
- LASSO
- Gaussian Processes
- Previous papers on this subject

Chapter 3

Results

Plan:

- Introduce a simple example function
- Possibly also use a more complex example?
- Training a neural network
  - Keras/Tensorflow
  - Choosing the right size NN
  - Order of the input points matters
- Fitting a linear regression
- Using stepwise regression
- Using LASSO
- Using GPs

To demonstrate the possibility of using statistical methods to understand the process of deep learning, we use a simple function $f(x) = x + 5\sin(x)$. 256 evenly spread datapoints were taken from this function, and noise following $\epsilon \sim \mathcal{N}(0, 0.1)$ was applied.

3.1 Training the neural network

This function was learnt by a neural network with 8 layers, each with 10 neurons and the $\tanh(\cdot)$ activation function, except for the final layer which used a linear activation function. The result of this learning is seen in Figure 3.1.

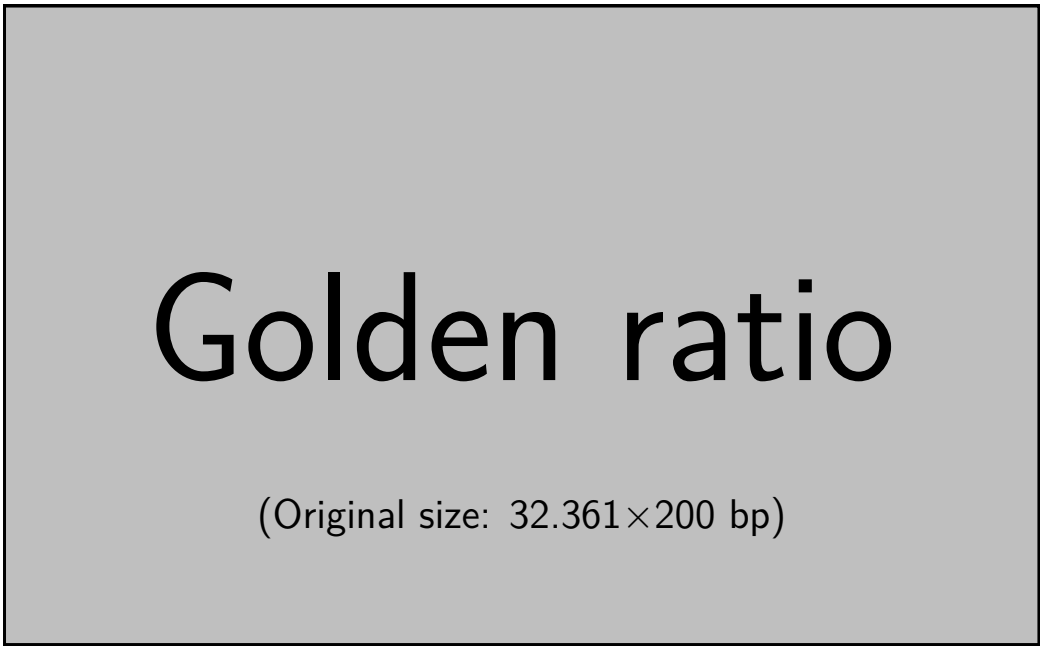

Figure 3.1: The output of the neural network.

3.1.1 Ordering of the datapoints

Due to the form of stochastic gradient descent used to train the model, if the order of the datapoints is not randomised, then the optimisation algorithm can more easily get stuck in a local minimum. An example of this is seen in Figure 3.2, where the same neural netowrk has been trained on the example data.

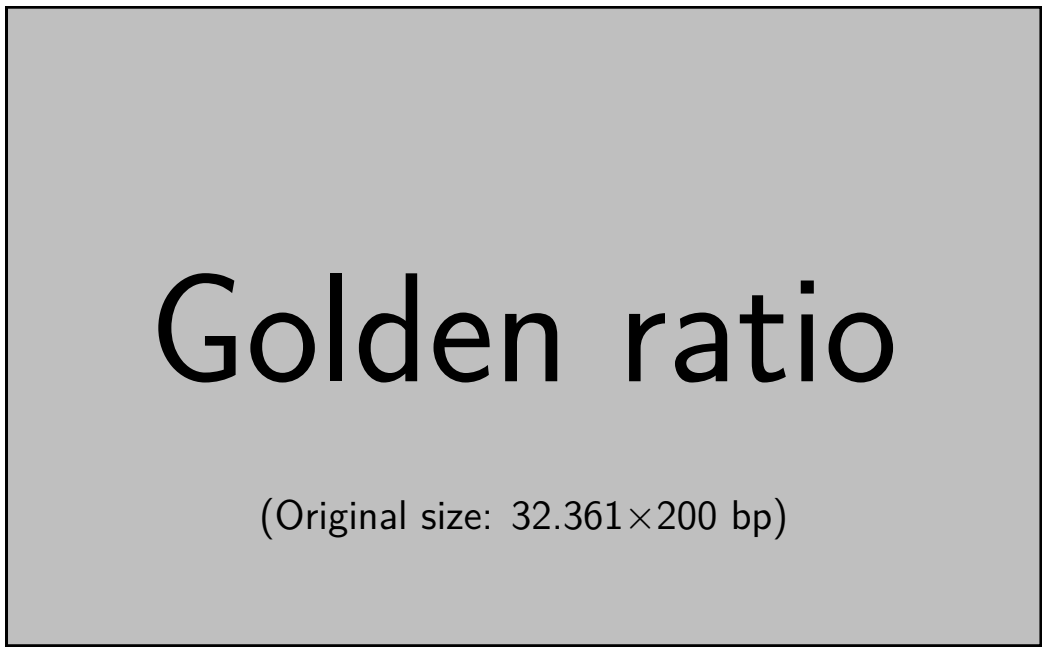

Figure 3.2: The output of the same neural network trained on the same data but ordered differently.

Chapter 4

Conclusion

This is chapter four.

- How well have each of the attempts worked?
- What could be improved?
- How useful would more research on this topic be?
- What should future research on this topic focus on?

Appendix A

Code

Appendix here.

Bibliography

Gandin, L. S. (1966). 'Objective analysis of meteorological fields'. Trans. by M. tirgume ha-mada' ha Yiśre'eli. *Quarterly Journal of the Royal Meteorological Society* 92.393, pp. 447–447. DOI: 10.1002/qj.49709239320.